

Лабораторная работа по Hive.

Выполнила Бабушкина Татьяна

1. Запустим контейнер с установленным hive и командой docker ps смотрим на наши контейнеры.

```
(base) tatyana@tatyana-Inspiron-5559:~/Desktop/lab2$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
6dd28edec209	cloudera/quickstart:latest	"/usr/bin/docker-qu..."	10 minutes ago	Up 10 minutes	0.0.0.0:8888->8888/tcp	compass
ionate_wing						
ed3f36f6a2e3	bde2020/hadoop-datanode:2.0.0-hadoop3.2.1-java8	"/entrypoint.sh /run..."	7 months ago	Up 43 minutes (healthy)	9864/tcp	datanod
e						
4c877fb07c44	bde2020/hadoop-nodemanager:2.0.0-hadoop3.2.1-java8	"/entrypoint.sh /run..."	7 months ago	Up 43 minutes (healthy)	8042/tcp	nodeman
ager						
93c33bd129ce	bde2020/hadoop-namenode:2.0.0-hadoop3.2.1-java8	"/entrypoint.sh /run..."	7 months ago	Up 43 minutes (healthy)	0.0.0.0:9000->9000/tcp, 0.0.0.0:9870->9870/tcp	namenod
e						
0b1d3b23e33a	bde2020/hadoop-resourcemanager:2.0.0-hadoop3.2.1-java8	"/entrypoint.sh /run..."	7 months ago	Up 41 minutes (healthy)	8088/tcp	resourc
emanager						
f2ac16187b00	bde2020/hadoop-historyserver:2.0.0-hadoop3.2.1-java8	"/entrypoint.sh /run..."	7 months ago	Up 43 minutes (healthy)	8188/tcp	history
server						

```
(base) tatyana@tatyana-Inspiron-5559:~/Desktop/lab2$ docker ps -n=10 --format '{{.ContainerID}} {{.Image}} {{.Command}} {{.Created}} {{.Status}} {{.Ports}} {{.Names}}'
```

2. Поместим данные в контейнер.

```
(base) tatyana@tatyana-Inspiron-5559:~/Desktop/lab2$ docker cp movies.dat 6dd28edec209:/volume/
(base) tatyana@tatyana-Inspiron-5559:~/Desktop/lab2$ docker cp ratings.dat 6dd28edec209:/volume/
(base) tatyana@tatyana-Inspiron-5559:~/Desktop/lab2$ docker cp users.dat 6dd28edec209:/volume/
```

3. Находим нужный контейнер (compassionate_wing) и переходим в него:

```
(base) tatyana@tatyana-Inspiron-5559:~/Desktop/lab2$ sudo docker exec -it compassionate_wing bash
[root@quickstart /]# ls
bin      etc      lib64    mnt      packer-files /sbin      sys      var
boot    home    lost+found  movies.t  proc          selinux   tmp      volume
dev      lib      media    opt      root          srv       usr
```

4. Конвертируем все данные.

```
[root@quickstart /]# cd volume
[root@quickstart volume]# ls
movies.dat ratings.dat users.dat
[root@quickstart volume]# sed 's/:/,/g' movies.dat > movies.t
[root@quickstart volume]# sed 's/:/,/g' ratings.dat > ratings.t
[root@quickstart volume]# sed 's/:/,/g' users.dat > users.t
[root@quickstart volume]# ls
movies.dat movies.t ratings.dat ratings.t users.dat users.t
```

5. Создаем директории под наши данные.

```
[root@quickstart volume]# hdfs dfs -mkdir -p /moviedata/ratings
[root@quickstart volume]# hdfs dfs -mkdir -p /moviedata/users
[root@quickstart volume]# hdfs dfs -mkdir -p /moviedata/movies
```

6. Переносим данные в hdfs.

```
[root@quickstart volume]# hdfs dfs -put movies.t /moviedata/movies
[root@quickstart volume]# hdfs dfs -put users.t /moviedata/users
[root@quickstart volume]# hdfs dfs -put ratings.t /moviedata/ratings
```

7. Используем Beeline:

```
[root@quickstart volume]# beeline -u jdbc:hive2://quickstart:10000/default -n admin -d org.apache.hive.jdbc.HiveDriver
2020-10-30 17:52:36,009 WARN [main] mapreduce.TableMapReduceUtil: The hbase-prefix-tree module jar containing PrefixTreeCodec is not present. Continuing without it.
Connecting to jdbc:hive2://quickstart:10000/default
Connected to: Apache Hive (version 1.1.0-cdh5.7.0)
Driver: Hive JDBC (version 1.1.0-cdh5.7.0)
Transaction isolation: TRANSACTION_REPEATABLE_READ
Beeline version 1.1.0-cdh5.7.0 by Apache Hive
```

8. Создаем БД.

```
Beeline version 1.1.0-cdh5.7.0 by Apache Hive
0: jdbc:hive2://quickstart:10000/default> create database moviedata;
INFO : Compiling command(queryId=hive_20201030175656_a0f9e3b9-f67d-4638-aa1a-05de80f55c4a): create database moviedata
INFO : Semantic Analysis Completed
INFO : Returning Hive schema: Schema(fieldSchemas:null, properties:null)
INFO : Completed compiling command(queryId=hive_20201030175656_a0f9e3b9-f67d-4638-aa1a-05de80f55c4a); Time taken: 0.736 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hive_20201030175656_a0f9e3b9-f67d-4638-aa1a-05de80f55c4a): create database moviedata
INFO : Starting task [Stage-0:DDL] in serial mode
INFO : Completed executing command(queryId=hive_20201030175656_a0f9e3b9-f67d-4638-aa1a-05de80f55c4a); Time taken: 0.65 seconds
INFO : OK
No rows affected (1.674 seconds)
0: jdbc:hive2://quickstart:10000/default> use moviedata;
INFO : Compiling command(queryId=hive_20201030175757_99caa96d-5962-4f3c-a78f-c858eaa60600): use moviedata
INFO : Semantic Analysis Completed
INFO : Returning Hive schema: Schema(fieldSchemas:null, properties:null)
INFO : Completed compiling command(queryId=hive_20201030175757_99caa96d-5962-4f3c-a78f-c858eaa60600); Time taken: 0.214 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hive_20201030175757_99caa96d-5962-4f3c-a78f-c858eaa60600): use moviedata
INFO : Starting task [Stage-0:DDL] in serial mode
INFO : Completed executing command(queryId=hive_20201030175757_99caa96d-5962-4f3c-a78f-c858eaa60600); Time taken: 0.006 seconds
INFO : OK
No rows affected (0.247 seconds)
```

```
0: jdbc:hive2://quickstart:10000/default> CREATE EXTERNAL TABLE moviedata (movieid INT, title STRING) ROW
```


10. Создаем таблицы для наших данных.

```
0: jdbc:hive2://quickstart:10000/default> CREATE EXTERNAL TABLE movies (userid INT, movie STRING, category STRING) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE LOCATION '/moviedata/movies';
INFO : Compiling command(queryId=hive_20200406112121_9c8ee76b-136a-4eab-ae2-355861407c85): CREATE EXTERNAL TABLE movies (userid INT, movie STRING, category STRING) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE LOCATION '/moviedata/movies'
INFO : Semantic Analysis Completed
INFO : Returning Hive schema: Schema(fieldSchemas:null, properties:null)
INFO : Completed compiling command(queryId=hive_20200406112121_9c8ee76b-136a-4eab-ae2-355861407c85); Time taken: 0.028 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hive_20200406112121_9c8ee76b-136a-4eab-ae2-355861407c85): CREATE EXTERNAL TABLE movies (userid INT, movie STRING, category STRING) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE LOCATION '/moviedata/movies'
INFO : Starting task [Stage-0:DDL] in serial mode
INFO : Completed executing command(queryId=hive_20200406112121_9c8ee76b-136a-4eab-ae2-355861407c85); Time taken: 0.077 seconds
INFO : OK
No rows affected (0.127 seconds)
0: jdbc:hive2://quickstart:10000/default> select * from movies limit 5;
INFO : Compiling command(queryId=hive_20200406112121_456dd395-f237-4a4b-ba6d-890fb834eb01): select * from movies limit 5
INFO : Semantic Analysis Completed
INFO : Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name:movies.userid, type:int, comment:null), FieldSchema(name:movies.movie, type:string, comment:null)], properties:null)
INFO : Completed compiling command(queryId=hive_20200406112121_456dd395-f237-4a4b-ba6d-890fb834eb01); Time taken: 0.136 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hive_20200406112121_456dd395-f237-4a4b-ba6d-890fb834eb01): select * from movies limit 5
INFO : Completed executing command(queryId=hive_20200406112121_456dd395-f237-4a4b-ba6d-890fb834eb01); Time taken: 0.0 seconds
INFO : OK
+-----+-----+-----+
| movies.userid | movies.movie | movies.category |
+-----+-----+-----+
| 1 | Toy Story (1995) | Animation|Children's|Comedy |
| 2 | Jumanji (1995) | Adventure|Children's|Fantasy |
| 3 | Grumpier Old Men (1995) | Comedy|Romance |
| 4 | Waiting to Exhale (1995) | Comedy|Drama |
+-----+-----+-----+

0: jdbc:hive2://quickstart:10000/default> CREATE EXTERNAL TABLE ratings (userid INT, movieId INT, rating INT, tstmp STRING) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE LOCATION '/moviedata/ratings';
INFO : Compiling command(queryId=hive_20200406111717_e4d97f6b-2310-4474-9bb0-dfc39a356f2c): CREATE EXTERNAL TABLE ratings (userid INT, movieId INT, rating INT, tstmp STRING) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE LOCATION '/moviedata/ratings'
INFO : Semantic Analysis Completed
INFO : Returning Hive schema: Schema(fieldSchemas:null, properties:null)
INFO : Completed compiling command(queryId=hive_20200406111717_e4d97f6b-2310-4474-9bb0-dfc39a356f2c); Time taken: 0.023 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hive_20200406111717_e4d97f6b-2310-4474-9bb0-dfc39a356f2c): CREATE EXTERNAL TABLE ratings (userid INT, movieId INT, rating INT, tstmp STRING) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE LOCATION '/moviedata/ratings'
INFO : Starting task [Stage-0:DDL] in serial mode
INFO : Completed executing command(queryId=hive_20200406111717_e4d97f6b-2310-4474-9bb0-dfc39a356f2c); Time taken: 0.767 seconds
INFO : OK
No rows affected (0.818 seconds)
0: jdbc:hive2://quickstart:10000/default> select * from ratings limit 5;
INFO : Compiling command(queryId=hive_20200406111717_0de20edf-9b37-4d1e-8ada-c670f7b31c4b): select * from ratings limit 5
INFO : Semantic Analysis Completed
INFO : Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name:ratings.userid, type:int, comment:null), FieldSchema(name:ratings.movieId, type:int, comment:null), FieldSchema(name:ratings.rating, type:int, comment:null), FieldSchema(name:ratings.tstmp, type:string, comment:null)], properties:null)
INFO : Completed compiling command(queryId=hive_20200406111717_0de20edf-9b37-4d1e-8ada-c670f7b31c4b); Time taken: 0.307 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hive_20200406111717_0de20edf-9b37-4d1e-8ada-c670f7b31c4b): select * from ratings limit 5
INFO : Completed executing command(queryId=hive_20200406111717_0de20edf-9b37-4d1e-8ada-c670f7b31c4b); Time taken: 0.001 seconds
INFO : OK
+-----+-----+-----+
| ratings.userid | ratings.movieId | ratings.rating | ratings.tstmp |
+-----+-----+-----+
| 1 | 1193 | 5 | 978300760 |
| 1 | 661 | 3 | 978302109 |
| 1 | 914 | 3 | 978301968 |
| 1 | 3408 | 4 | 978300275 |
| 1 | 2355 | 5 | 978824291 |
+-----+-----+-----+
5 rows selected (0.562 seconds)
```

11. Пишем сами запросы:

```
0: jdbc:hive2://quickstart:10000/default> select * from movies m where m.category = 'Comedy' limit 10;
INFO : Compiling command(queryId=hive_20200406120606_f9ddcd7a-01ed-4735-a7b7-be487719ed73): select * from movies m where m.category = 'Comedy' limit 10
INFO : Semantic Analysis Completed
INFO : Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name:m.userid, type:int, comment:null), FieldSchema(name:m.movie, type:string, comment:null), FieldSchema(name:m.category, type:string, comment:null)], properties:null)
INFO : Completed compiling command(queryId=hive_20200406120606_f9ddcd7a-01ed-4735-a7b7-be487719ed73); Time taken: 0.311 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hive_20200406120606_f9ddcd7a-01ed-4735-a7b7-be487719ed73): select * from movies m where m.category = 'Comedy' limit 10
INFO : Completed executing command(queryId=hive_20200406120606_f9ddcd7a-01ed-4735-a7b7-be487719ed73); Time taken: 0.0 seconds
INFO : OK

+-----+-----+-----+
| m.userid | m.movie | m.category |
+-----+-----+-----+
| 5 | Father of the Bride Part II (1995) | Comedy |
| 19 | Ace Ventura: When Nature Calls (1995) | Comedy |
| 38 | It Takes Two (1995) | Comedy |
| 52 | Mighty Aphrodite (1995) | Comedy |
| 63 | Don't Be a Menace to South Central While Drinking Your Juice in the Hood (1996) | Comedy |
| 65 | Bio-Dome (1996) | Comedy |
| 69 | Friday (1995) | Comedy |
| 88 | Black Sheep (1996) | Comedy |
| 96 | In the Bleak Midwinter (1995) | Comedy |
| 101 | Bottle Rocket (1996) | Comedy |
+-----+-----+-----+
10 rows selected (0.404 seconds)
```



```

0: jdbc:hive2://quickstart:10000/default> select ratings.movieid,max(ratings.rating) from ratings group by ratings.movieid limit 10;
INFO : Compiling command(queryId=hive_20200406120303_c0a24f6e-f2d5-4e92-a905-0a839e747f99): select ratings.movieid,max(ratings.rating) from ratings group by ratings.movieid limit 10
INFO : Semantic Analysis Completed
INFO : Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name:ratings.movieid, type:int, comment:null), FieldSchema(name:_c1, type:int, comment:null)], properties:null)
INFO : Completed compiling command(queryId=hive_20200406120303_c0a24f6e-f2d5-4e92-a905-0a839e747f99); Time taken: 0.161 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hive_20200406120303_c0a24f6e-f2d5-4e92-a905-0a839e747f99): select ratings.movieid,max(ratings.rating) from ratings group by ratings.movieid limit 10
INFO : Query ID = hive_20200406120303_c0a24f6e-f2d5-4e92-a905-0a839e747f99
INFO : Total jobs = 1
INFO : Launching Job 1 out of 1
INFO : Starting task [Stage-1:MAPRED] in serial mode
INFO : Number of reduce tasks not specified. Estimated from input data size: 1
INFO : In order to change the average load for a reducer (in bytes):
INFO :   set hive.exec.reducers.bytes.per.reducer=<number>
INFO : In order to limit the maximum number of reducers:
INFO :   set hive.exec.reducers.max=<number>
INFO : In order to set a constant number of reducers:
INFO :   set mapreduce.job.reduces=<number>
INFO : Starting Job = job_1586167015334_0008, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1586167015334_0008/
INFO : Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1586167015334_0008
INFO : Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
INFO : 2020-04-06 12:03:26,650 Stage-1 map = 0%,   reduce = 0%
INFO : 2020-04-06 12:03:51,326 Stage-1 map = 100%,   reduce = 0%, Cumulative CPU 13.98 sec
INFO : 2020-04-06 12:04:52,204 Stage-1 map = 100%,   reduce = 0%, Cumulative CPU 13.98 sec
INFO : 2020-04-06 12:04:54,281 Stage-1 map = 100%,   reduce = 100%, Cumulative CPU 16.74 sec
INFO : MapReduce Total cumulative CPU time: 16 seconds 740 msec
INFO : Ended Job = job_1586167015334_0008
INFO : MapReduce Jobs Launched:
INFO : Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 16.74 sec HDFS Read: 21600885 HDFS Write: 41 SUCCESS
INFO : Total MapReduce CPU Time Spent: 16 seconds 740 msec
INFO : Completed executing command(queryId=hive_20200406120303_c0a24f6e-f2d5-4e92-a905-0a839e747f99); Time taken: 97.459 seconds
INFO : OK

```

```

+-----+-----+
| ratings.movieid | _c1 |
+-----+-----+
| 1               | 5   |
| 2               | 5   |
| 3               | 5   |
| 4               | 5   |
| 5               | 5   |
| 6               | 5   |
| 7               | 5   |
| 8               | 5   |
| 9               | 5   |
| 10              | 5   |
+-----+-----+
10 rows selected (97.674 seconds)

```

```

0: jdbc:hive2://quickstart:10000/default> select movies.category, count(movies.category) from movies group by movies.category limit 10;
INFO : Compiling command(queryId=hive_20200406114747_da588977-86c8-416f-a220-b9e31c867610): select movies.category, count(movies.category) from movies group by movies.category limit 10
INFO : Semantic Analysis Completed
INFO : Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name=movies.category, type:string, comment:null), FieldSchema(name=_c1, type:bigint, comment:null)], properties:null)
INFO : Completed compiling command(queryId=hive_20200406114747_da588977-86c8-416f-a220-b9e31c867610); Time taken: 0.283 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hive_20200406114747_da588977-86c8-416f-a220-b9e31c867610): select movies.category, count(movies.category) from movies group by movies.category limit 10
INFO : Query ID = hive_20200406114747_da588977-86c8-416f-a220-b9e31c867610
INFO : Total jobs = 1
INFO : Launching Job 1 out of 1
INFO : Starting task [Stage-1:MAPRED] in serial mode
INFO : Number of reduce tasks not specified. Estimated from input data size: 1
INFO : In order to change the average load for a reducer (in bytes):
INFO :   set hive.exec.reducers.bytes.per.reducer=<number>
INFO : In order to limit the maximum number of reducers:
INFO :   set hive.exec.reducers.max=<number>
INFO : In order to set a constant number of reducers:
INFO :   set mapreduce.job.reduces=<number>
INFO : Starting Job = job_1586167015334_0004, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1586167015334_0004/
INFO : Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1586167015334_0004
INFO : Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
INFO : 2020-04-06 11:47:20,210 Stage-1 map = 0%, reduce = 0%
INFO : 2020-04-06 11:47:45,935 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 8.55 sec
INFO : 2020-04-06 11:48:36,405 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 10.89 sec
INFO : 2020-04-06 11:48:38,564 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 12.0 sec
INFO : MapReduce Total cumulative CPU time: 12 seconds 0 msec
INFO : Ended Job = job_1586167015334_0004
INFO : MapReduce Jobs Launched:
INFO : Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 12.0 sec HDFS Read: 170861 HDFS Write: 161 SUCCESS
INFO : Total MapReduce CPU Time Spent: 12 seconds 0 msec
INFO : Completed executing command(queryId=hive_20200406114747_da588977-86c8-416f-a220-b9e31c867610); Time taken: 91.583 seconds
INFO : OK

```

movies.category	_c1
20th Century (Zbogum na dvadesetiot vek) (1998)	1
302 (1995)	1
A (1920)	1
A (1923)	1
A (1932)	1
A (1937)	2
A (1938)	1
A (1945)	1
A (1951)	2
A (1957)	2

10 rows selected (91.93 seconds)

```

0: jdbc:hive2://quickstart:10000/default> select count(*) from users where users.age > 50;
INFO : Compiling command(queryId=hive_20200406115353_07e7439e-fe5b-427c-8918-56661aaf269a): select count(*) from users where users.age > 50
INFO : Semantic Analysis Completed
INFO : Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name: c0, type:bigint, comment:null)], properties:null)
INFO : Completed compiling command(queryId=hive_20200406115353_07e7439e-fe5b-427c-8918-56661aaf269a); Time taken: 0.211 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hive_20200406115353_07e7439e-fe5b-427c-8918-56661aaf269a): select count(*) from users where users.age > 50
INFO : Query ID = hive_20200406115353_07e7439e-fe5b-427c-8918-56661aaf269a
INFO : Total jobs = 1
INFO : Launching Job 1 out of 1
INFO : Starting task [Stage-1:MAPRED] in serial mode
INFO : Number of reduce tasks determined at compile time: 1
INFO : In order to change the average load for a reducer (in bytes):
INFO :   set hive.exec.reducers.bytes.per.reducer=<number>
INFO : In order to limit the maximum number of reducers:
INFO :   set hive.exec.reducers.max=<number>
INFO : In order to set a constant number of reducers:
INFO :   set mapreduce.job.reduces=<number>
INFO : Starting Job = job_1586167015334_0006, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1586167015334_0006/
INFO : Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1586167015334_0006
INFO : Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
INFO : 2020-04-06 11:53:41,311 Stage-1 map = 0%, reduce = 0%
INFO : 2020-04-06 11:53:55,603 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.3 sec
INFO : 2020-04-06 11:54:04,381 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.49 sec
INFO : MapReduce Total cumulative CPU time: 4 seconds 490 msec
INFO : Ended Job = job_1586167015334_0006
INFO : MapReduce Jobs Launched:
INFO : Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.49 sec HDFS Read: 117617 HDFS Write: 4 SUCCESS
INFO : Total MapReduce CPU Time Spent: 4 seconds 490 msec
INFO : Completed executing command(queryId=hive_20200406115353_07e7439e-fe5b-427c-8918-56661aaf269a); Time taken: 32.973 seconds
INFO : OK
+-----+
| _c0 |
+-----+
| 380 |
+-----+
1 row selected (33.239 seconds)

```



```

1 row selected (33.123 seconds)
0: jdbc:hive2://quickstart:10000/default> select u.*,m.* from users u inner join movies m on u.movieid=m.movieid limit 10 ;
Error: Error while compiling statement: FAILED: SemanticException [Error 10002]: Line 1:53 Invalid column reference 'movieid' (state=42000,code=10002)
0: jdbc:hive2://quickstart:10000/default> select u.*,m.* from users u inner join movies m on u.movieid = m.movieid limit 10 ;
Error: Error while compiling statement: FAILED: SemanticException [Error 10002]: Line 1:53 Invalid column reference 'movieid' (state=42000,code=10002)
0: jdbc:hive2://quickstart:10000/default> select u.*,m.* from users u inner join movies m on u.userid = m.userid limit 10 ;
INFO : Compiling command(queryId=hive_20200406120000_45fa77e2-1093-4946-b2a3-54ac76e41461): select u.*,m.* from users u inner join movies m on u.userid = m.userid limit 10
INFO : Semantic Analysis Completed
INFO : Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name:u.userid, type:int, comment:null), FieldSchema(name:u.gender, type:string, comment:null), FieldSchema(name:u.age, type:int, comment:null), FieldSchema(name:u.unk1, type:int, comment:null), FieldSchema(name:u.unk2, type:int, comment:null), FieldSchema(name:m.userid, type:int, comment:null), FieldSchema(name:m.movie, type:string, comment:null)], FieldSchema(name:m.category, type:string, comment:null]), properties:null)
INFO : Completed compiling command(queryId=hive_20200406120000_45fa77e2-1093-4946-b2a3-54ac76e41461); Time taken: 0.211 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hive_20200406120000_45fa77e2-1093-4946-b2a3-54ac76e41461): select u.*,m.* from users u inner join movies m on u.userid = m.userid limit 10
INFO : Query ID = hive_20200406120000_45fa77e2-1093-4946-b2a3-54ac76e41461
INFO : Total jobs = 1
INFO : Starting task [Stage-4:MAPREDLOCAL] in serial mode
INFO : Execution completed successfully
INFO : MapredLocal task succeeded
INFO : Launching Job 1 out of 1
INFO : Starting task [Stage-3:MAPRED] in serial mode
INFO : Number of reduce tasks is set to 0 since there's no reduce operator
INFO : Starting Job = job_1586167015334_0007, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1586167015334_0007/
INFO : Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1586167015334_0007
INFO : Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
INFO : 2020-04-06 12:00:34,677 Stage-3 map = 0%, reduce = 0%
INFO : 2020-04-06 12:00:52,338 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 14.04 sec
INFO : MapReduce Total cumulative CPU time: 14 seconds 40 msec
INFO : Ended Job = job_1586167015334_0007
INFO : MapReduce Jobs Launched:
INFO : Stage-Stage-3: Map: 1 Cumulative CPU: 14.04 sec HDFS Read: 10611 HDFS Write: 558 SUCCESS
INFO : Total MapReduce CPU Time Spent: 14 seconds 40 msec
INFO : Completed executing command(queryId=hive_20200406120000_45fa77e2-1093-4946-b2a3-54ac76e41461); Time taken: 37.59 seconds
INFO : OK
+-----+-----+-----+-----+-----+-----+-----+-----+
| u.userid | u.gender | u.age | u.unk1 | u.unk2 | m.userid | m.movie | m.category |
+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | F | 1 | 10 | 48067 | 1 | Toy Story (1995) | Animation|Children's|Comedy |
| 2 | M | 56 | 16 | 70072 | 2 | Jumanji (1995) | Adventure|Children's|Fantasy |
| 3 | M | 25 | 15 | 55117 | 3 | Grumpier Old Men (1995) | Comedy|Romance |
| 4 | M | 45 | 7 | 2460 | 4 | Waiting to Exhale (1995) | Comedy|Drama |
| 5 | M | 25 | 20 | 55455 | 5 | Father of the Bride Part II (1995) | Comedy |
| 6 | F | 50 | 9 | 55117 | 6 | Heat (1995) | Action|Crime|Thriller |
| 7 | M | 35 | 1 | 6810 | 7 | Sabrina (1995) | Comedy|Romance |
| 8 | M | 25 | 12 | 11413 | 8 | Tom and Huck (1995) | Adventure|Children's |
| 9 | M | 25 | 17 | 61614 | 9 | Sudden Death (1995) | Action |
| 10 | F | 35 | 1 | 95370 | 10 | GoldenEye (1995) | Action|Adventure|Thriller |
+-----+-----+-----+-----+-----+-----+-----+-----+
10 rows selected (38.01 seconds)

```

:)